

Memformer

A Memory Guided Transformer for Time Series Forecasting

Yunyao Cheng, Chenjuan Guo, Bin Yang, Haomin Yu, Kai Zhao, Christian S. Jensen

February 2025

Proceedings of the VLDB Endowment, Volume 18, Issue 2

Presented by **Andreas Gottschalk Krath**

1. Introduction



1.1 Motivation

Forecasting

- Predicting the future
 - Allows preparation
- Many applications
 - Electricity prices
 - Finance
- Long term forecasting?
 - Obviously more difficult than short term
 - Time constrained tasks

1.1 Motivation

Long Term Forecasting

- What defines long term?
 - Historical horizon
 - Forecasting horizon
 - Both exceed 96 time steps
 - Hourly time step \rightarrow 4 days
 - Time series

Variable Correlation

- Complex systems have many variables
- A increases and B increases \rightarrow Positive
- A increases and B decreases \rightarrow Negative
- A increases and B is stagnant \rightarrow None
- These impact forecasting accuracy
 - Patterns in the data

1.1 Motivation

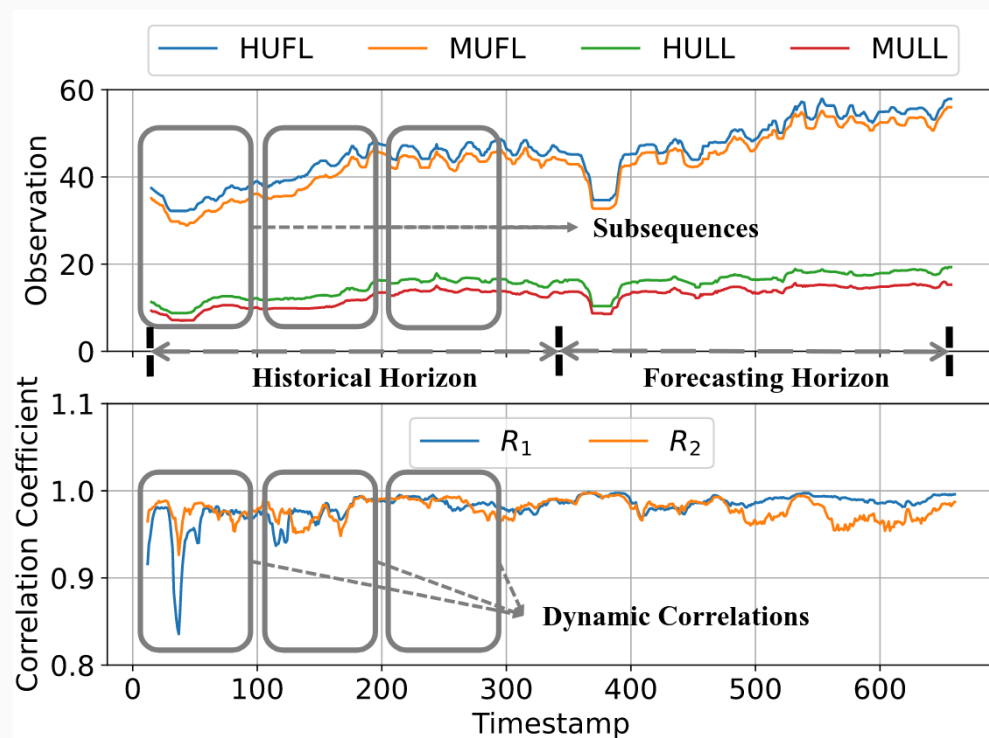
Dynamic Correlations

- Are variable correlations stable over time?
 - No
- Correlations are dynamic over time
 - Seasons
 - Sensor drift
- We often consider average
 - Especially hurtful in time series
 - Predictions are bad in periods

1.1 Motivation

Dynamic Correlations

- Are variable correlations stable over time?
 - No
- Correlations are dynamic over time
 - Seasons
 - Sensor drift
- We often consider average
 - Especially hurtful in time series
 - Predictions are bad in periods



(a) Dynamic correlations. The Average $R_1 = 0.995$ and $R_2 = 0.990$.

1.1 Motivation


Disrupted Correlations

- System errors
- External influence
- What happens with outliers?
 - Affect correlation \rightarrow accuracy
- Many models are sensitive to outliers
 - Numeric difference dominates training
 - Reason for a lot of preprocessing
 - Normalization
 - Clipping
 - Pruning

1.1 Motivation

Disrupted Correlations

- System errors
- External influence
- What happens with outliers?
 - Affect correlation \rightarrow accuracy
- Many models are sensitive to outliers
 - Numeric difference dominates training
 - Reason for a lot of preprocessing
 - Normalization
 - Clipping
 - Pruning



(b) Disrupted correlation. The Average $R_1 = 0.908$ and $R_2 = 0.963$.

1.2 Problem

Challenge 1

- Capture dynamic correlations
- Mitigate disrupted correlations
- Existing solutions struggle with the latter
 - Capture dynamic and disrupted
 - Reduces model robustness

Challenge 2

- Local information 🤝 global information
- Global information is *all* local information
- Local information *affects* global information
- Existing solutions struggle with combining
 - Only local
 - Only global

1.3 Contributions

Memformer

- Transformer
- Patch-wise recurrent graph learning
 - Captures dynamic correlations
- Global attention
 - Mitigates disrupted correlations
- Addresses challenge 1

Alternating Memory Enhancer

- Memory network
- Associates local and global information
- Addresses challenge 2

Experiments

- Proof

2. Methodology

2.2 Preprocessing

Instance normalization

- Normalize within historical horizon only
- Mitigates the issue of internal covariate shift
- Allows model to effectively grasp the intricate temporal dynamics inherent in time series

$$H' = (H - \mu) / \sqrt{(\sigma^2 + c)}, \text{ where}$$

H is the historical horizon

μ is the mean

σ is the variance

c ensures numerical stability

2.2 Preprocessing

Instance normalization

- Normalize within historical horizon only
- Mitigates the issue of internal covariate shift
- Allows model to effectively grasp the intricate temporal dynamics inherent in time series

$$H' = (H - \mu) / \sqrt{(\sigma^2 + c)}, \text{ where}$$

H is the historical horizon

μ is the mean

σ is the variance

c ensures numerical stability

poral dynamics inherent in time series. Instance normalization is defined as $\mathbf{H}' = (\mathbf{H} - \mu) / \sqrt{(\sigma^2 + \text{constant})}$, where \mathbf{H}' denotes the preprocessed feature, μ and σ denote the mean and variance of the sample, respectively, and “constant” is a small positive real number included to ensure numerical stability.

2.2 Preprocessing

Instance normalization

- Normalize within historical horizon only
- Mitigates the issue of internal covariate shift
- Allows model to effectively grasp the intricate temporal dynamics inherent in time series

$$H' = (H - \mu) / \sqrt{(\sigma^2 + c)}, \text{ where}$$

H is the historical horizon

μ is the mean

σ is the variance

c ensures numerical stability

poral dynamics inherent in time series. Instance normalization is defined as $\mathbf{H}' = (\mathbf{H} - \mu) / \sqrt{(\sigma^2 + \text{constant})}$, where \mathbf{H}' denotes the preprocessed feature, μ and σ denote the mean and variance of the sample, respectively, and “constant” is a small positive real number included to ensure numerical stability.

- Mistake in variance?
 - σ is conventional notation for standard deviation
 - σ^2 is conventional notation for variance

2.2 Preprocessing

What is going on?

2.2 Preprocessing

What is going on?

- Explored code to find answer
- `data_provider/data_loader.py`
 - Only place anything related to loading data happens
 - `Dataset_ETT_hour`, `Dataset_ETT_minute`, `Dataset_Custom`, `Dataset_Pred`

2.2 Preprocessing

What is going on?

- Explored code to find answer
- data_provider/data_loader.py
 - Only place anything related to loading data happens
 - Dataset_ETT_hour, Dataset_ETT_minute, Dataset_Custom, Dataset_Pred

```
from sklearn.preprocessing import StandardScaler
class ...:
    def __read_data__(self):
        self.scaler = StandardScaler()
        self.scaler.fit(train_data.values)
        data = self.scaler.transform(df_data.values)
```

2.2 Preprocessing

What is going on?

- Explored code to find answer
- `data_provider/data_loader.py`
 - Only place anything related to loading data happens
 - `Dataset_ETT_hour`, `Dataset_ETT_minute`, `Dataset_Custom`, `Dataset_Pred`

```
from sklearn.preprocessing import StandardScaler
class ...:
    def __read_data__(self):
        self.scaler = StandardScaler()
        self.scaler.fit(train_data.values)
        data = self.scaler.transform(df_data.values)
```

- They fit on training data
- Normalize entire dataset with μ and σ from training data

2.2 Preprocessing

What are they actually doing?

Preprocessing

$$H' = (H - \mu) / \sqrt{(\sigma^2 + c)}, \text{ where}$$

H is the historical horizon

μ is the mean

σ is the variance

c ensures numerical stability

StandardScaler

$$z = (x - \mu) / \sigma, \text{ where}$$

x is the sample

μ is the mean

σ is the standard deviation

2.2 Preprocessing

What are they actually doing?

Preprocessing

$$H' = (H - \mu) / \sqrt{(\sigma^2 + c)}, \text{ where}$$

H is the historical horizon

μ is the mean

σ is the variance

c ensures numerical stability

- We know that $\sqrt{\sigma^2} = \sigma$

StandardScaler

$$z = (x - \mu) / \sigma, \text{ where}$$

x is the sample

μ is the mean

σ is the standard deviation

2.2 Preprocessing

What are they actually doing?

Preprocessing

$$H' = (H - \mu) / \sqrt{(\sigma^2 + c)}, \text{ where}$$

H is the historical horizon

μ is the mean

σ is the variance

c ensures numerical stability

- We know that $\sqrt{\sigma^2} = \sigma$
- Essentially same formula, except constant

StandardScaler

$$z = (x - \mu) / \sigma, \text{ where}$$

x is the sample

μ is the mean

σ is the standard deviation

2.2 Preprocessing

What are they actually doing?

Preprocessing

$$H' = (H - \mu) / \sqrt{(\sigma^2 + c)}, \text{ where}$$

H is the historical horizon

μ is the mean

σ is the variance

c ensures numerical stability

- We know that $\sqrt{\sigma^2} = \sigma$
- Essentially same formula, except constant
- Fit on training data, normalize entire dataset \rightarrow global normalization

StandardScaler

$$z = (x - \mu) / \sigma, \text{ where}$$

x is the sample

μ is the mean

σ is the standard deviation

2.2 Preprocessing

What are they actually doing?

Preprocessing

$$H' = (H - \mu) / \sqrt{(\sigma^2 + c)}, \text{ where}$$

H is the historical horizon

μ is the mean

σ is the variance

c ensures numerical stability

StandardScaler

$$z = (x - \mu) / \sigma, \text{ where}$$

x is the sample

μ is the mean

σ is the standard deviation

- We know that $\sqrt{\sigma^2} = \sigma$
- Essentially same formula, except constant
- Fit on training data, normalize entire dataset \rightarrow global normalization
- None of the stated benefits of instance normalization
 - Mitigate internal covariate shift
 - Grasp intricate temporal dynamics in TS

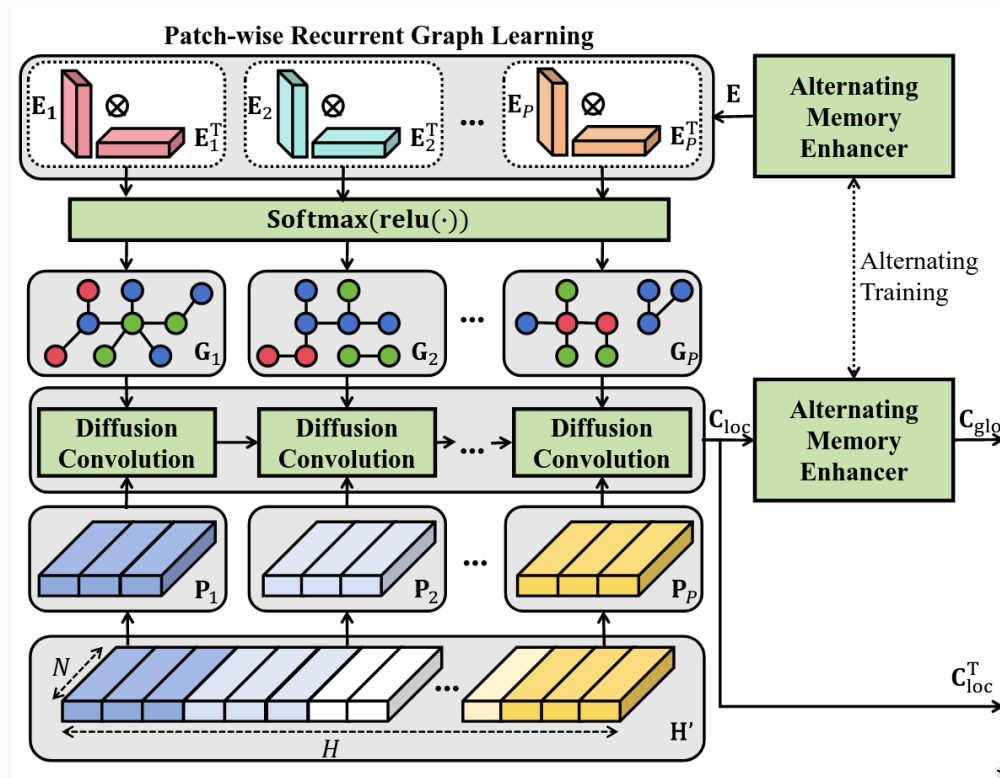
2.3 Patch-wise Recurrent Graph Learning

Architecture

Upper part \rightarrow dynamic correlation

Lower part \rightarrow normalized data

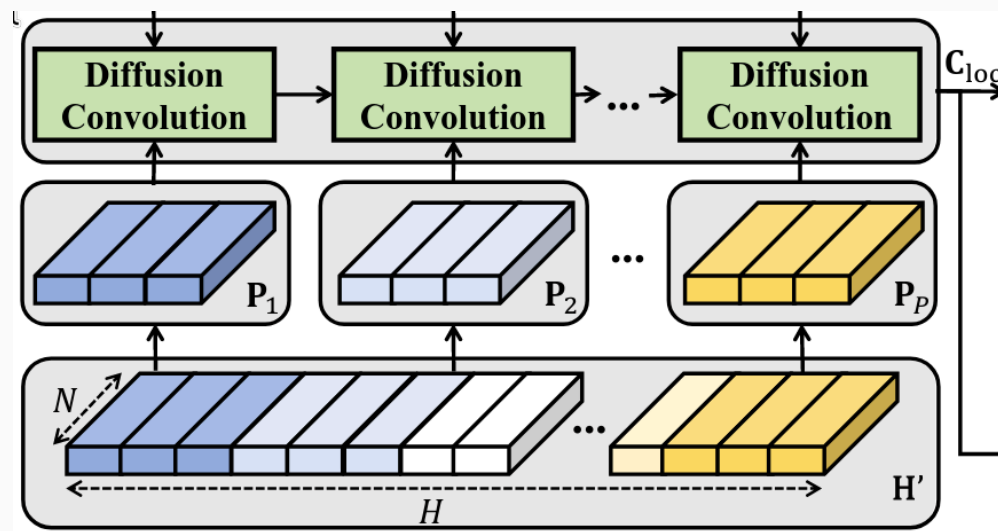
Output \rightarrow enriched input features



2.3 Patch-wise Recurrent Graph Learning

Normalized Data

- Normalized as described earlier
 - Not what the paper actually states



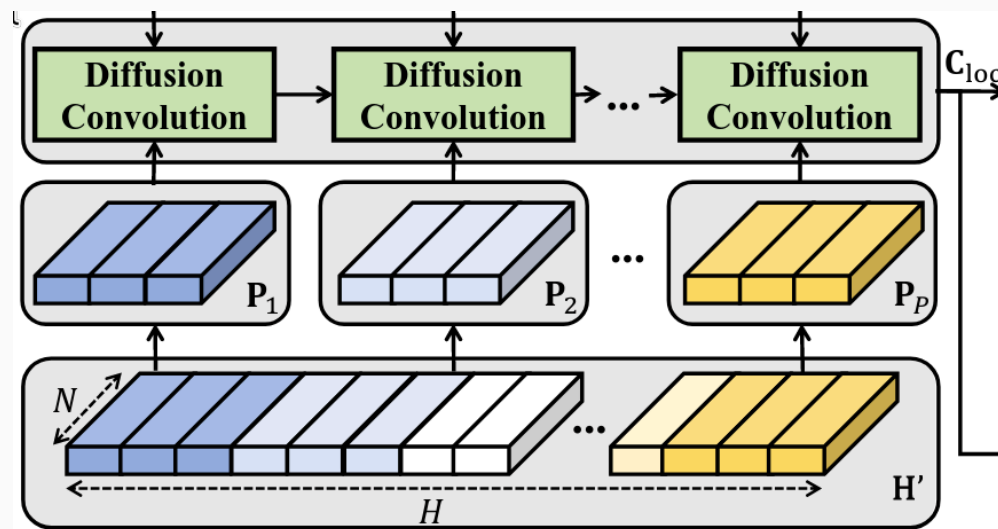
2.3 Patch-wise Recurrent Graph Learning

Normalized Data

- Normalized as described earlier
 - Not what the paper actually states

Patches

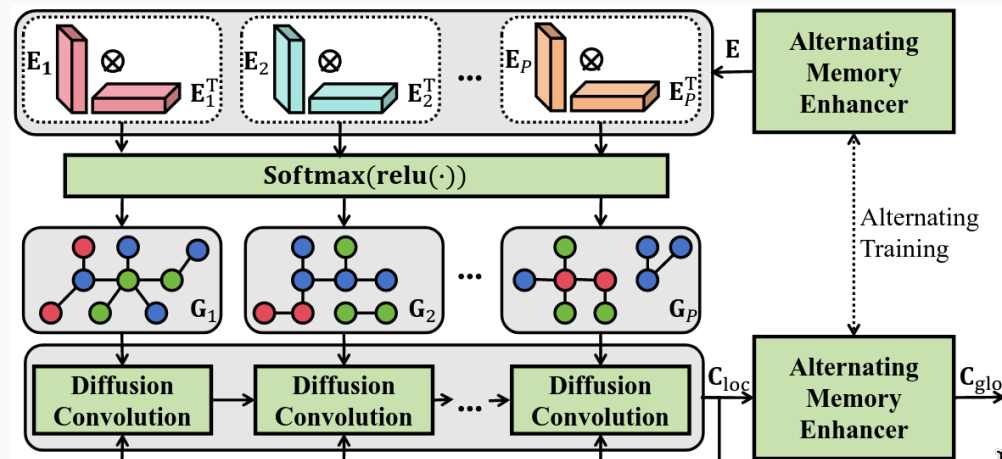
- H' is split into p patches
- Stride S
- Size T
- If $S \geq T$ patches are disjoint
- If $S < T$ patches overlap
 - Common elements for adjacent patches



2.3 Patch-wise Recurrent Graph Learning

AME

- Provides local memory embedding
 - These are learnable parameters
- Consistent local memory for patch P_i
- Matrix product of $E_i \otimes E_i^T$
 - Similarity matrix for variables in P_i




2.3 Patch-wise Recurrent Graph Learning

AME

- Provides local memory embedding
 - These are learnable parameters
- Consistent local memory for patch P_i
- Matrix product of $E_i \otimes E_i^T$
 - Similarity matrix for variables in P_i

ReLU + Softmax

- ReLU eliminates negative values
 - Removes negative correlations
- Softmax scales into influence scores



2.3 Patch-wise Recurrent Graph Learning

AME

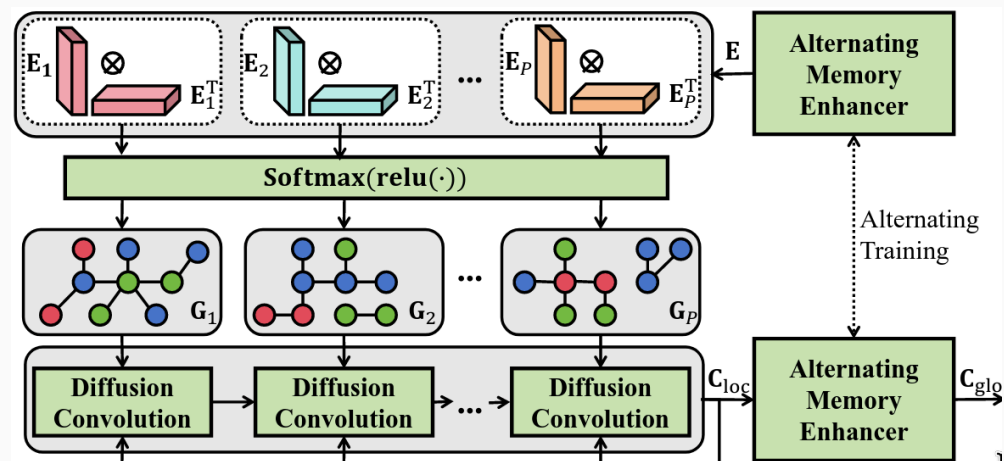
- Provides local memory embedding
 - These are learnable parameters
- Consistent local memory for patch P_i
- Matrix product of $E_i \otimes E_i^T$
 - Similarity matrix for variables in P_i

ReLU + Softmax

- ReLU eliminates negative values
 - Removes negative correlations
- Softmax scales into influence scores

Graph

- Translates influence scores into graph
- Captures connection between variables
 - Dynamic correlations



2.3 Patch-wise Recurrent Graph Learning

Diffusion Convolution


- Normalized data is adjusted based on connections in graph
- Numeric values “diffuse” into neighbours
 - Not only immediate neighbours
- Spatially relates data based on connections

Gated Recurrent Unit

- Forwards information from P_i to P_{i+1}
- Temporally relates data in a sequence

Output

- Input features enriched with local information
- Spatial \rightarrow dynamic correlations
- Temporal \rightarrow GRU



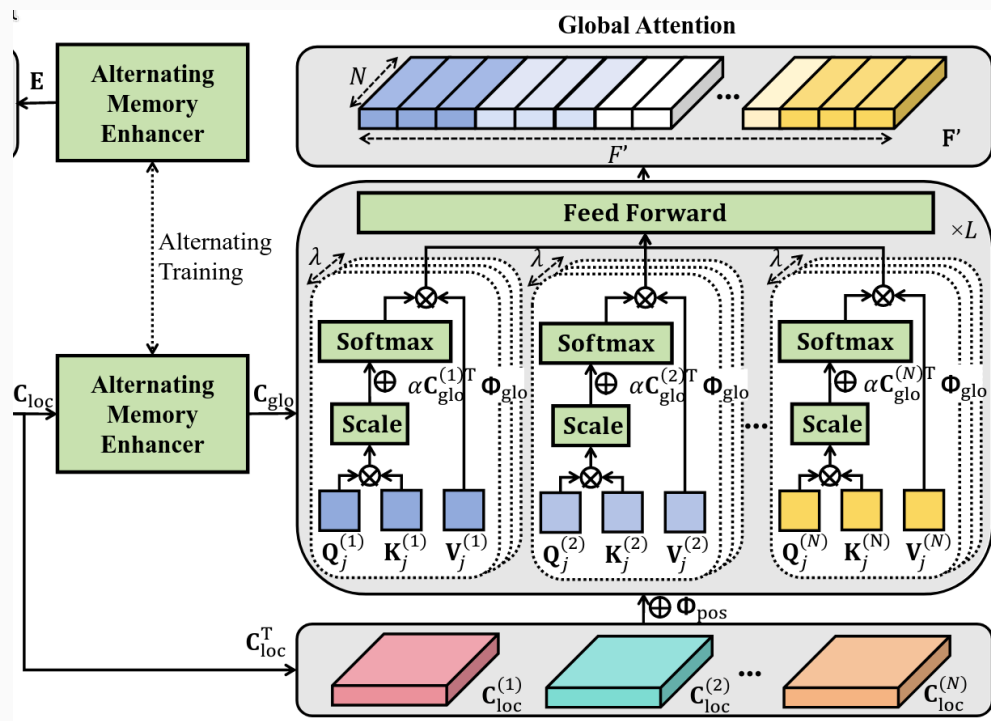
2.4 Global Attention

Motivation

- Patch-wise correlations are sensitive
 - Outliers dominate
- Constrain locally enriched features
 - Mitigate disrupted correlations

Input

- Transpose locally enriched features
 - Isolate variables
 - Diffusion earlier
- Linear transformation
 - Positional encoding
- Converted to Q, K, V matrices
 - Learnable parameters




2.4 Global Attention

Attention

- Relatively conventional implementation
 - Query and Key to find importance
 - Weight Value by importance
- Global information is new
- Adding global information after softmax
 - Bias probabilities
 - Global information affects parameters

Output

- The final “representation” of data
- F' is not a forecast
 - Final feature representation
- Linear layer maps to forecasting horizon




2.5 Alternating Memory Enhancer

Overview

- Input
 - Locally correlated features
- Outputs
 - Local information E
 - Global information C_{glo}
- Shared global memory

Hyperparameters


- $M \rightarrow$ number of high level patterns
 - Spikes, seasons, stable
- $D_{\text{glo}} \rightarrow$ richness of patterns



2.5 Alternating Memory Enhancer

Local Enhancer


- Local memory regions Γ_{loci}
 - One for each patch
- $P_i \longleftrightarrow \Gamma_{loci} \longrightarrow \Gamma_{loci} \longleftrightarrow E_i$
- Not directly identical
 - E contains global memories
 - Defined by C_{loc}
- Memories are **not** information



2.5 Alternating Memory Enhancer

Global Enhancer

- Learns from locally correlated features
- Γ_{mem} is a large trainable tensor
 - Produces inquiry tensor
 - Recognizes patterns in data
 - The M high level patterns
- Inquiry tensor
 - Prevalence of patterns in local data
 - Similarity scores with global memory
- Probability distribution
 - Importance of pattern
- Top k most important patterns
 - Stored in C_{glo}
 - Scaled based on importance
 - Weighted sum



2.5 Alternating Memory Enhancer

Alternating Training

- Local information E requires
 - Local memories
 - Global memories
- Updating both simultaneously
 - Unstable training
 - Issues converging
- LE and GE alternate training
 - Split adjustment of memories

LE Training

- Local memories > global memories
 - More parameters → longer convergence
- Balance convergence
 - Different learning rates
 - LE training more

2.5 Alternating Memory Enhancer

Alternating Training

- Local information E requires
 - Local memories
 - Global memories
- Updating both simultaneously
 - Unstable training
 - Issues converging
- LE and GE alternate training
 - Split adjustment of memories

LE Training

- Local memories > global memories
 - More parameters \rightarrow longer convergence
- Balance convergence
 - Different learning rates
 - LE training more

Algorithm 1 AME alternating training

Input: Historical horizon and ground truth \mathbf{H}, \mathbf{F} ; local and global memories $\Gamma_{\text{loc}}, \Gamma_{\text{glo}}$; local training step ϵ ; learning rates $\eta_{\text{loc}}, \eta_{\text{glo}}$ for local and global enhancers

Output: Local and global information $\mathbf{E}, \mathbf{C}_{\text{glo}}$; learned local and global memories $\Gamma_{\text{loc}}, \Gamma_{\text{glo}}$, tensor Γ_{mem} , and bias \mathbf{b}_{mem}

- 1: *Initialisation:* Initializing local and global memories $\Gamma_{\text{loc}}, \Gamma_{\text{glo}}$, tensor Γ_{mem} , and bias \mathbf{b}_{mem} randomly
- 2: **while** $\Gamma_{\text{loc}}, \Gamma_{\text{glo}}, \Gamma_{\text{mem}}$, and \mathbf{b}_{mem} are not converged **do**
- 3: **for** iteration = 0 to ϵ **do**
- 4: $\mathbf{H}' \leftarrow \text{Preprocessing}(\mathbf{H})$
- 5: $\mathbf{E} \leftarrow \mathcal{A}_{\text{loc}}(\Gamma_{\text{loc}}, \Gamma_{\text{glo}})$
- 6: $\mathbf{C}_{\text{loc}} \leftarrow \mathcal{G}_{\Theta}(\mathbf{H}', \mathbf{E})$
- 7: $\mathbf{C}_{\text{glo}} \leftarrow \mathcal{A}_{\text{glo}}(\mathbf{C}_{\text{loc}}, \Gamma_{\text{glo}})$
- 8: $\mathbf{F}' \leftarrow \mathcal{T}_{\Phi}(\mathbf{C}_{\text{loc}}, \mathbf{C}_{\text{glo}})$
- 9: $\hat{\mathbf{F}} \leftarrow \text{LinearHead}(\mathbf{F}')$
- 10: $\Gamma_{\text{loc}} \leftarrow \Gamma_{\text{loc}} - \eta_{\text{loc}} \nabla_{\Gamma_{\text{loc}}} \mathcal{L}(\hat{\mathbf{F}}, \mathbf{F})$
- 11: **end for**
- 12: $\mathbf{H}' \leftarrow \text{Preprocessing}(\mathbf{H})$
- 13: $\mathbf{E} \leftarrow \mathcal{A}_{\text{loc}}(\Gamma_{\text{loc}}, \Gamma_{\text{glo}})$
- 14: $\mathbf{C}_{\text{loc}} \leftarrow \mathcal{G}_{\Theta}(\mathbf{H}', \mathbf{E})$
- 15: $\mathbf{C}_{\text{glo}} \leftarrow \mathcal{A}_{\text{glo}}(\mathbf{C}_{\text{loc}}, \Gamma_{\text{glo}})$
- 16: $\mathbf{F}' \leftarrow \mathcal{T}_{\Phi}(\mathbf{C}_{\text{loc}}, \mathbf{C}_{\text{glo}})$
- 17: $\hat{\mathbf{F}} \leftarrow \text{LinearHead}(\mathbf{F}')$
- 18: $\Gamma_{\text{glo}} \leftarrow \Gamma_{\text{glo}} - \eta_{\text{glo}} \nabla_{\Gamma_{\text{glo}}} \mathcal{L}(\hat{\mathbf{F}}, \mathbf{F})$
- 19: $\Gamma_{\text{mem}} \leftarrow \Gamma_{\text{mem}} - \eta_{\text{glo}} \nabla_{\Gamma_{\text{mem}}} \mathcal{L}(\hat{\mathbf{F}}, \mathbf{F})$,
- 20: $\mathbf{b}_{\text{mem}} \leftarrow \mathbf{b}_{\text{mem}} - \eta_{\text{glo}} \nabla_{\mathbf{b}_{\text{mem}}} \mathcal{L}(\hat{\mathbf{F}}, \mathbf{F})$
- 21: **end while**

3. Results

4. Critique

4.1 Notation

Preprocessing

- As mentioned earlier
- Unconventional notation
- Obscures details

Inconsistencies

- C_{glo} is global memory
- C_{loc} is locally correlated features
- E is local memory

4.1 Notation

Preprocessing

- As mentioned earlier
- Unconventional notation
- Obscures details

Inconsistencies

- C_{glo} is global memory
- C_{loc} is locally correlated features
- E is local memory

Symbol Reuse

- \mathbf{F} is the ground truth
- F is the dimensionality of \mathbf{F}
- \mathbf{F}' is the encoding output
- F' is the dimensionality of \mathbf{F}'
- Confusing statements and diagrams

4.1 Notation

Preprocessing

- As mentioned earlier
- Unconventional notation
- Obscures details

Inconsistencies

- C_{glo} is global memory
- C_{loc} is locally correlated features
- E is local memory

Symbol Reuse

- \mathbf{F} is the ground truth
- F is the dimensionality of \mathbf{F}
- \mathbf{F}' is the encoding output
- F' is the dimensionality of \mathbf{F}'
- Confusing statements and diagrams

$\mathbf{F}' \in \mathbb{R}^{F' \times N}$, where F' is the temporal dimension of the representation.

4.1 Notation

Preprocessing

- As mentioned earlier
- Unconventional notation
- Obscures details



Inconsistencies

- C_{glo} is global memory
- C_{loc} is locally correlated features
- E is local memory

Symbol Reuse

- \mathbf{F} is the ground truth
- F is the dimensionality of \mathbf{F}
- \mathbf{F}' is the encoding output
- F' is the dimensionality of \mathbf{F}'
- Confusing statements and diagrams

$\mathbf{F}' \in \mathbb{R}^{F' \times N}$, where F' is the temporal dimension of the representation.



5. Praise

5.1 Figures

Colors

- Help understanding and data flow
 - Preprocessing → final encoding
 - Minor inconsistencies
 - Attention

Dimensionality

- Squares → 2-dimensional
- Cubes → 3-dimensional
- Transposed → lying down
- Slices of shapes
 - M slices of global memory
 - P slices of local memory

